

Lecture 1: Laying the Foundations + Terminology

Chapters 1.1-1.2

Goals for Today

- ▶ Go over the syllabus
- ▶ Show some fun examples
- ▶ Discuss how to evaluate the efficacy of a **treatment**
- ▶ Describe the different kinds of **variables** we'll consider

What is statistics?

(Direct from text) The general scientific process of investigation can be summed up as follows:


1. Identify the scientific question or problem
2. Collect relevant data on the topic
3. Analyze the data
4. Form a conclusion and communicate it

Statistics concerns itself with points 2 through 4.


Your Majors

Biology	Biochem and Molecular Biology
8	4
Psychology	Economics
3	2
Environmental Studies	Biology/Psychology
2	1
English	Environmental Studies-Biol
1	1
Environmental Studies-Hist	Physics
1	1
Political Science	
1	

Example: 2012 Election - Nate Silver's Predictions vs Actual Results



Nate Silver's Map



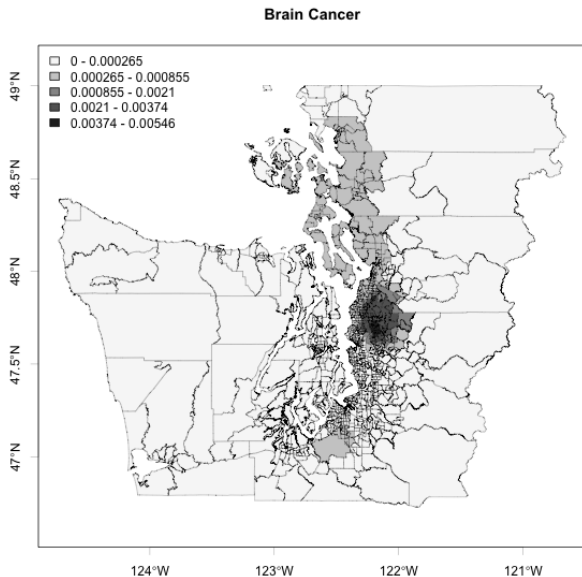
The Actual Map

Example: Brain & Breast Cancer in Western Washington


My PhD dissertation involved detecting cancer “clusters”: areas of residual spatial variation of disease risk.

We modeled the (Bayesian) probability of cluster membership for each of the $n = 887$ census tracts in Western Washington in 2000, using cancer data from 1995–2005, controlling for age, race, and gender.


Brain Cancer Controlling for Age, Race, & Gender




Breast Cancer Controlling for Age, Race, & Gender




Income per Capita Quintiles



Breast Cancer Adjusted for Income as Well



Example: Social Network Display of a Recent Party I Had



Math Colloquium

- ▶ What: “Mathematical Models for HIV Transmission Networks”
- ▶ Who: UW Statistics and Sociology Prof. Martina Morris (Reed '80 sociology alum)
- ▶ Where: Physics 123
- ▶ When: Thursday January 29th

Say we want answer the following questions:

- ▶ Does a new kind of cognitive therapy alter levels of depression in patients?
- ▶ Or you question the effectiveness of antioxidants in preventing cancer.
- ▶ Will reassuring potential new users to a gambling website that we won't spam them increase the sign-up rate?

Evaluating the efficacy of a 'treatment'

In all the above cases, you are questioning the efficacy of a **treatment/intervention**. One way to evaluate the efficacy is via an **experiment** where you define

- ▶ A **control** group: the “business as usual” baseline group
- ▶ A **treatment** group: the group that receives/is subject to the treatment/intervention

and make comparisons.

Website Experiments

Control:

Join BettingExpert

Username:

Email:

Password:

☐ I accept the [Terms and Conditions](#)

Sign up +



Treatment:

Join BettingExpert

Username:

Email:

Password:

☐ I accept the [Terms and Conditions](#)
100% privacy - we will never spam you!

Sign up +

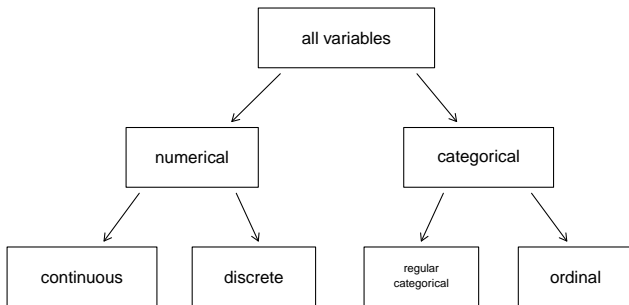
Example of a treatment vs control

Two other examples in the media of late

- ▶ Facebook's tinkering with user's emotions ([link](#))
- ▶ OkCupid's admission that they experiment on human beings ([link](#))

Variables

A **variable** is a description of any characteristic whose value may change from one unit in the population to the next:



Variables

A **variable** is a description of any characteristic whose value may change from one unit in the population to the next:

1. gender of an engineer: **categorical** variable
2. level of education (high school/GED, college, grad school): **ordinal** variable
3. number of major defects on a newly manufactured phone: **discrete** variable i.e. something you can count
4. temperature of the battery in a phone after 1 hour of use: **continuous** variable; its possible values consist of an interval on the number line

Data

At its simplest, data are presented in a data table or matrix where (almost always) each

- ▶ row corresponds to **cases** or **units of observation/analysis**
- ▶ column represents the variables corresponding to a particular observation

It is almost always the case that

- ▶ n is the number of observations
- ▶ p is the number of variables

Data Summaries

Consider the variable "federal spending per capita" in each of the 3,143 counties in the US. One can hardly digest this:

[1]	6.068095	6.139862	8.752158	7.122016	5.130910	9.973062	9.311835	15.439218
[9]	8.613707	7.104621	6.324061	10.640378	9.781442	8.982702	6.840035	20.330684
[17]	9.687698	11.080738	7.839761	9.461856	9.650295	7.760627	25.774791	13.948106
...								
[3121]	7.520731	10.246400	3.106800	17.679572	4.824044	7.247212	8.484211	8.794626
[3129]	9.829593	8.100945	17.090715	4.855849	6.621378	22.587359	10.813260	11.422522
[3137]	9.580265	4.368986	5.062138	6.236968	4.549105	8.713817	6.694784	

Data Summaries


We can't interpret all the data at once; we need to boil it down via [summary statistics](#), single numbers summarizing a large amount of data.

Using the `summary()` command in R:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
0.000	6.964	8.669	9.991	10.860	204.600	4

Relationships between variables

We can best display the relationship between two variables using a scatterplot AKA bivariate plot:



Relationships between variables

Almost always we are interested in the relationship between two or more variables.

A pair of variables are either related in some way (**associated**) or not (**independent**). No pair of variables are both associated and independent.

We can have either a **negative association** (as the value of one variable increases, the other decreases) or a **positive association**.

Relationships between variables

We can consider a third variable in the previous plot.

