

Lecture 6: Visualizing Numerical and Categorical Data

Chapter 1.6+1.7

1 / 25

Goals for Today

- ▶ Rule of thumb for standard deviations
- ▶ Population vs sample mean/variance/standard deviations
- ▶ Percentiles and Quartiles
- ▶ Boxplots
- ▶ Piecharts, barplots, mosaicplots

2 / 25


Rule of Thumb for Standard Deviations

If the data distribution is bell-shaped, then

- ▶ about $\frac{2}{3}$ of the data will be within one SD of the mean (book says 70%).
- ▶ about 95% of the data will be within two SD.

3 / 25

Example




- ▶ black line is mean \bar{x}
- ▶ red lines mark about $\frac{2}{3}$:
 $[\bar{x} - s, \bar{x} + s] =$
 $[20 - 3, 20 + 3] = [17, 23].$
- ▶ blue lines mark about 95%:
 $[\bar{x} - 2s, \bar{x} + 2s] =$
 $[20 - 6, 20 + 6] = [14, 26].$

4 / 25

Population vs Sample Mean/Variance/Standard Deviation

Recall the notion of taking a **representative sample** from a **study/target population**. Say we are interested in the income of the individuals.



5 / 25

Population vs Sample Mean/Variance/Standard Deviation

- ▶ The **sample mean** \bar{x} is the mean income of the 4 sampled people.
- ▶ The **population mean** μ is the mean income of all 24 people in the target population.
- ▶ We say \bar{x} **estimates** μ . If the sample is representative, then \bar{x} estimates μ with high **accuracy** i.e. it is unbiased.

6 / 25

Population vs Sample Mean/Variance/Standard Deviation

	True Population Value	Sample Value
Mean	μ	\bar{x}
Variance	σ^2	s^2
Standard Deviation	σ	s

The sample value is used to **estimate** the (true) population value.

7 / 25

Percentiles

A percentile (%'ile) indicates the value **below** which a given %'age of observations fall.

SAT Scores from 2012

<http://media.collegeboard.com/digitalServices/pdf/research/SAT-Percentile-Ranks-2012.pdf>

So for example, if you scored 700 in critical reading, 95% of college-bound seniors who took the test did worse.

8 / 25

Quartiles and IQR

Quartiles split up the data into 4 intervals, each with about one quarter of the data:

- ▶ The lower quartile is the 25th %'ile
- ▶ The median is the 50th %'ile
- ▶ The upper quartile is the 75th %'ile


The **interquartile range (IQR)** is another measure of the spread of a sample:

$$\text{IQR} = \text{upper quartile} - \text{lower quartile}$$

9 / 25

MLB Data Quartiles

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
400.0	418.3	1094.0	3282.0	4250.0	33000.0



The IQR is (3rd Quartile - 1st Quartile) = $4250.0 - 418.3 = 3831.7$
i.e the distance between the red and blue line.

10 / 25

Robust Statistics (Chapter 1.6.6)

Robust estimates are statistics where extreme observations (outliers) have less effect on their values, i.e. are more resistant to their effect. The median and IQR are two examples.

Example: Old scoring system in figure skating: **drop the highest & lowest scores** and then take the average.

Say we have a figure skater who gets judged by countries V-Z:

Country	V	W	X	Y	Z
Score	4.0	5.2	5.2	5.3	6.0

Drop the 4.0 and 6.0, then the final score is: $\frac{5.2+5.2+5.3}{3} = 5.23$

11 / 25

Boxplots

Boxplots are visual summaries of a sample x_1, \dots, x_n that bring to light unusual values (potential outliers):


Example: # US Forces casualties in the war in Afghanistan for each month from 2008-2009:

7, 1, 7, 5, 16, 28, 20, 22, 27, 16, 1, 3, 14, 15, 13, 6, 12, 24, 44, 51, 37, 59, 17, 17

12 / 25

Boxplots

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.00	7.00	16.00	19.25	24.75	59.00



US Forces casualties in Afghanistan for each month 2008-2009

Page 29 of text describes the length of the **whiskers**: they capture data that is no more than $1.5 \times IQR$ of both ends of the box.

13 / 25

Outliers Are Relatively Extreme

An **outlier** is an observation that appears extreme relative to the rest of the data.


Why it is important to look for outliers? Examination of data for possible outliers serves many useful purposes, including

- ▶ Identifying strong skew in the distribution.
- ▶ Identifying data collection or entry errors.
- ▶ Providing insight into interesting properties of the data.

14 / 25

Piecharts

Say we have the following piecharts represent the polling from a local election with five candidates (1-5) at three different time points A, B, and C:




Answer the following questions:

- ▶ In the first race, is candidate 5 doing better than candidate 4?
- ▶ Who did better between time A and time B, candidate 2 or candidate 4?

15 / 25

Barplots Instead




Answers:

- ▶ Candidate 5 is doing better than 4
- ▶ Between A and B, candidate 2 went from about 17% to 20% while candidate 4 went from about 22% to 21%. So candidate 2 did better

16 / 25

3D Piecharts Can Be Deceiving



EPP (teal) has 266 seats, whereas S&D (red) has 190 seats.

17 / 25

Titanic Survival Data

Typing `data(Titanic)` in R loads the survival and death counts, split by each of the following categories:

- ▶ Class: 1st, 2nd, 3rd, or crew (4 levels)
- ▶ Gender (2 levels)
- ▶ Age: Child or adult (2 levels)

i.e. $4 \times 2 \times 2 = 16$ possible groups to consider.

Questions

- ▶ What was the effect of class (1st, 2nd, 3rd, crew) on your chances of survival?
- ▶ Did the “women and children” first lifeboat policy hold?

18 / 25

Frequency Table


A table summarizing a single categorical variable is called a **frequency table**. Overall:

Died	1490
Survived	711
Total	2201

19 / 25

Barplot

Barplots are ways to display categorical variables:



20 / 25

Contingency Table

A table that **cross-classifies** two categorical variables is a **contingency table**. Now let's split survival by class: 1st, 2nd, 3rd, and crew.

Before:

Died	1490
Survived	711
Total	2201


After:

	1st	2nd	3rd	Crew	Total
Died	122	167	528	673	1490
Survived	203	118	178	212	711
Total	325	285	706	885	2201

21 / 25

Stacked Barplot


Stacked barplots are one way to display values from a contingency table:



22 / 25

Mosaic Plots


Mosaic plots are similar, but the widths of the bars now reflect proportions:



23 / 25

Stacked Barplots

Using the `ggplot2` package, we can plot survivals by class, age, and gender all at once.



24 / 25

Standardized/Normalized Stacked Barplots

Instead of raw counts, we can expand each bar to reflect proportions (i.e. standardize/normalize them).

