

Lecture 17: Paired Data and Difference of Two Means

Chapter 5.2, 5.1

Goals for Today

- ▶ Difference of means
- ▶ Paired differences of means
- ▶ Note on Practical vs Statistical Significance

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?
4. What is the proportion p of “successes”?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?
4. What is the proportion p of “successes”?
5. Are the proportions of “successes” of two groups p_1 and p_2 equal or not?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?
4. What is the proportion p of “successes”?
5. Are the proportions of “successes” of two groups p_1 and p_2 equal or not?
6. Are the means μ_1, \dots, μ_k of k groups **all** equal or not?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:

1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?
4. What is the proportion p of “successes”?
5. Are the proportions of “successes” of two groups p_1 and p_2 equal or not?
6. Are the means μ_1, \dots, μ_k of k groups **all** equal or not?
7. Are we observing counts that we were expected?

The 8 Types of Questions

Here are the 8 broad types of questions we can answer with statistical methods (confidence intervals and hypothesis tests) in this class:


1. What is the mean value μ ?
2. Are the means of two groups μ_1 and μ_2 equal or not?
3. What is the mean paired difference μ_{diff} ?
4. What is the proportion p of “successes”?
5. Are the proportions of “successes” of two groups p_1 and p_2 equal or not?
6. Are the means μ_1, \dots, μ_k of k groups **all** equal or not?
7. Are we observing counts that we were expected?
8. Are two categorical variables independent?

Chapter 5.2: Are the means of 2 groups μ_1 and μ_2 equal?

Did men (n=45) run faster than women (n=55)?

Chapter 5.2: Are the means of 2 groups μ_1 and μ_2 equal?

Did men (n=45) run faster than women (n=55)?



Difference in Means

We are interested in the difference of two population means:

Difference in Means

We are interested in the difference of two population means:

- ▶ μ_w is the mean time for women
- ▶ μ_m is the mean time for men

Difference in Means

We are interested in the difference of two population means:

- ▶ μ_w is the mean time for women
- ▶ μ_m is the mean time for men

The data:

	men	women
\bar{x}	87.65	102.13
s	12.5	15.2
n	45	55

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

1. Determine a point estimate of $\mu_w - \mu_m$.

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

1. Determine a point estimate of $\mu_w - \mu_m$.
2. Show the normality of the sampling distribution: mean and SE

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

1. Determine a point estimate of $\mu_w - \mu_m$.
2. Show the normality of the sampling distribution: mean and SE
3. Build a confidence interval

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

1. Determine a point estimate of $\mu_w - \mu_m$.
2. Show the normality of the sampling distribution: mean and SE
3. Build a confidence interval
4. Conduct hypothesis tests

Difference of Means

We now recreate all the elements of Chapter 4 using this new population parameter $\mu_w - \mu_m$:

1. Determine a point estimate of $\mu_w - \mu_m$.
2. Show the normality of the sampling distribution: mean and SE
3. Build a confidence interval
4. Conduct hypothesis tests

First, the point estimate for $\mu_w - \mu_m$ is the sample difference of means

$$\bar{x}_w - \bar{x}_m = 102.13 - 87.65 = 14.48$$

Normality of Sampling Distribution

If the sample means \bar{x}_1 and \bar{x}_2

Normality of Sampling Distribution

If the sample means \bar{x}_1 and \bar{x}_2

- ▶ each meet the criteria for having nearly normal sampling distributions

Normality of Sampling Distribution

If the sample means \bar{x}_1 and \bar{x}_2

- ▶ each meet the criteria for having nearly normal sampling distributions
- ▶ also the observations from the two samples are independent

Normality of Sampling Distribution

If the sample means \bar{x}_1 and \bar{x}_2

- ▶ each meet the criteria for having nearly normal sampling distributions
- ▶ also the observations from the two samples are independent

then the difference in sample means $\bar{x}_1 - \bar{x}_2$ will also have a nearly normal sampling distribution...

Normality of Sampling Distribution

with

Normality of Sampling Distribution

with

- ▶ mean $\mu_1 - \mu_2$

Normality of Sampling Distribution

with

- ▶ mean $\mu_1 - \mu_2$
- ▶ estimated standard error

$$SE_{\bar{x}_1 - \bar{x}_2} = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

Normality of Sampling Distribution

with

- ▶ mean $\mu_1 - \mu_2$
- ▶ estimated standard error

$$SE_{\bar{x}_1 - \bar{x}_2} = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

Note the different s^2 's and sample sizes.

Normality of Sampling Distribution

We verify the conditions:

Normality of Sampling Distribution

We verify the conditions:

- ▶ Because each sample consists of less than 10% of their respective populations (men: 45 of 7192 and women: 55 of 9732).

Normality of Sampling Distribution

We verify the conditions:

- ▶ Because each sample consists of less than 10% of their respective populations (men: 45 of 7192 and women: 55 of 9732).
- ▶ The observations for both groups don't look too skewed.

Normality of Sampling Distribution

We verify the conditions:

- ▶ Because each sample consists of less than 10% of their respective populations (men: 45 of 7192 and women: 55 of 9732).
- ▶ The observations for both groups don't look too skewed.
- ▶ Each sample has at least 30 observations (rule of thumb).

Normality of Sampling Distribution

We verify the conditions:

- ▶ Because each sample consists of less than 10% of their respective populations (men: 45 of 7192 and women: 55 of 9732).
- ▶ The observations for both groups don't look too skewed.
- ▶ Each sample has at least 30 observations (rule of thumb).
- ▶ The samples are independent (not paired or linked in any way).

Normality of Sampling Distribution

We verify the conditions:

- ▶ Because each sample consists of less than 10% of their respective populations (men: 45 of 7192 and women: 55 of 9732).
- ▶ The observations for both groups don't look too skewed.
- ▶ Each sample has at least 30 observations (rule of thumb).
- ▶ The samples are independent (not paired or linked in any way).

the sampling distribution is Normal with mean $= \mu_w - \mu_m$ and

$$SE_{\bar{x}_w - \bar{x}_m} = \sqrt{\frac{15.2^2}{55} + \frac{12.5^2}{45}} = 2.77$$

Confidence Interval

A 95% confidence interval for $\mu_1 - \mu_2$ is

$$(\text{point estimate for } \mu_1 - \mu_2) \pm 1.96 \times SE$$

Confidence Interval

A 95% confidence interval for $\mu_1 - \mu_2$ is

$$\begin{aligned} &(\text{point estimate for } \mu_1 - \mu_2) \pm 1.96 \times SE \\ &(\bar{x}_1 - \bar{x}_2) \pm 1.96 \times SE_{\bar{x}_1 - \bar{x}_2} \end{aligned}$$

Confidence Interval

A 95% confidence interval for $\mu_1 - \mu_2$ is

$$\begin{aligned} &(\text{point estimate for } \mu_1 - \mu_2) \pm 1.96 \times SE \\ &(\bar{x}_1 - \bar{x}_2) \pm 1.96 \times SE_{\bar{x}_1 - \bar{x}_2} \end{aligned}$$

So for the Cherry Blossom Run data, a 95% CI for $\mu_w - \mu_m$ is:

$$14.48 \pm 1.96 \times 2.77 = [9.05, 19.91]$$

Chapter 5.1: Paired Data

Two sets of observations are **paired** if each observation in one set has a special correspondence or connection with exactly one observation in the other data set.

Chapter 5.1: Paired Data

Two sets of observations are **paired** if each observation in one set has a special correspondence or connection with exactly one observation in the other data set.

Examples:

- ▶ Cholesterol levels before and after some intervention

Chapter 5.1: Paired Data

Two sets of observations are **paired** if each observation in one set has a special correspondence or connection with exactly one observation in the other data set.

Examples:

- ▶ Cholesterol levels before and after some intervention
- ▶ Disease rates amongst pairs of twins

Chapter 5.1: Paired Data


Two sets of observations are **paired** if each observation in one set has a special correspondence or connection with exactly one observation in the other data set.

Examples:

- ▶ Cholesterol levels before and after some intervention
- ▶ Disease rates amongst pairs of twins
- ▶ In the text: price of the same textbook at the UCLA bookstore vs Amazon


Paired Differences

The methodology for paired data remains the same, except our **observations** are the difference in pairs. Example, for the UCLA Bookstore vs Amazon book price example in the text



Paired Differences

The methodology for paired data remains the same, except our **observations** are the difference in pairs. Example, for the UCLA Bookstore vs Amazon book price example in the text



Paired Differences

We have

Paired Differences

We have

- ▶ population parameter is μ_{diff}

Paired Differences

We have

- ▶ population parameter is μ_{diff}
- ▶ point estimate \bar{x}_{diff} of μ_{diff}

Paired Differences

We have

- ▶ population parameter is μ_{diff}
- ▶ point estimate \bar{x}_{diff} of μ_{diff}
- ▶ Conditions: not on the original observations, but rather the **differences**: 10% rule, sample size n , and not too skewed *differences*.

Paired Differences

We have

- ▶ population parameter is μ_{diff}
- ▶ point estimate \bar{x}_{diff} of μ_{diff}
- ▶ Conditions: not on the original observations, but rather the **differences**: 10% rule, sample size n , and not too skewed *differences*.
- ▶ If met, \bar{x}_{diff} has a normal sampling distribution with mean μ_{diff} and $SE_{diff} = \frac{s_{diff}}{\sqrt{n_{diff}}}$.

Next Time

- ▶ One sample t-test